

```
In [48]: # import required libraries
import numpy as np
import pandas as pd
from sklearn.decomposition import PCA
```

```
In [49]: #import data set
mer_train= pd.read_csv('train.csv')
```

```
In [50]: # view the dataset first 10 rows
mer_train.head(10)
```

```
Out[50]:
```

	ID	y	X0	X1	X2	X3	X4	X5	X6	X8	...	X375	X376	X377	X378	X379	X380	X382	X
0	0	130.81	k	v	at	a	d	u	j	o	...	0	0	1	0	0	0	0	
1	6	88.53	k	t	av	e	d	y	l	o	...	1	0	0	0	0	0	0	
2	7	76.26	az	w	n	c	d	x	j	x	...	0	0	0	0	0	0	1	
3	9	80.62	az	t	n	f	d	x	l	e	...	0	0	0	0	0	0	0	
4	13	78.02	az	v	n	f	d	h	d	n	...	0	0	0	0	0	0	0	
5	18	92.93	t	b	e	c	d	g	h	s	...	0	0	1	0	0	0	0	
6	24	128.76	al	r	e	f	d	f	h	s	...	0	0	0	0	0	0	0	
7	25	91.91	o	l	as	f	d	f	j	a	...	0	0	0	0	0	0	0	
8	27	108.67	w	s	as	e	d	f	i	h	...	1	0	0	0	0	0	0	
9	30	126.99	j	b	aq	c	d	f	a	e	...	0	0	1	0	0	0	0	

10 rows × 378 columns



```
In [51]: # view the info of data
mer_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4209 entries, 0 to 4208
Columns: 378 entries, ID to X385
dtypes: float64(1), int64(369), object(8)
memory usage: 12.1+ MB
```

```
In [52]: # check for null values
mer_train.isnull().sum()
```

```
Out[52]: ID      0
         y      0
         X0     0
         X1     0
         X2     0
         ..
         X380   0
         X382   0
         X383   0
         X384   0
         X385   0
         Length: 378, dtype: int64
```

```
In [53]: # convert the 'y' values into array format
         # seperate the 'y' column from the dataset
y_train = mer_train['y'].values
```

```
In [54]: y_train
```

```
Out[54]: array([130.81,  88.53,  76.26, ..., 109.22,  87.48, 110.85])
```

```
In [55]: # count of features with having X in the name
cols= [c for c in mer_train.columns if 'X' in c]
print ('Number of features: {}'.format(len(cols)))
```

```
Number of features: 376
```

```
In [56]: # counts of each columns
print ('feature Types:')
mer_train[cols].dtypes.value_counts()
```

```
feature Types:
```

```
Out[56]: int64      368
         object      8
         dtype: int64
```

```
In [57]: # values count of constant features, binary features and Categorical features
counts = [[],[],[[]]
```

```
In [58]: for c in cols:
        typ = mer_train[c].dtype
        unique = len(np.unique(mer_train[c]))
        if unique == 1:
            counts[0].append(c)
        elif unique == 2 and typ == np.int64:
            counts[1].append(c)
        else:
            counts[2].append(c)

        print('Constant features: {} Binary features: {} Categorical features: {}'.format(*[len(c) for c in counts]))
    print('Constant features:', counts[0])
    print('Categorical features:', counts[2])
```

Constant features: 12 Binary features: 356 Categorical features: 8

Constant features: ['X11', 'X93', 'X107', 'X233', 'X235', 'X268', 'X289', 'X290', 'X293', 'X297', 'X330', 'X347']  
 Categorical features: ['X0', 'X1', 'X2', 'X3', 'X4', 'X5', 'X6', 'X8']

```
In [59]: # Load the test dataset
mer_test= pd.read_csv('test.csv')
```

```
In [60]: # select the usable columns
usable_columns = list(set(mer_train.columns)- set(['ID','y']))
y_train = mer_train['y'].values
id_test = mer_test['ID'].values

x_train = mer_train[usable_columns]
x_test = mer_test[usable_columns]
```

```
In [61]: # check for null values in the dataset
def check_null(df):
    if df.isnull().any().any():
        print('There are missing values in the dataset')
    else:
        print('There are no missing values in the dataset')

check_null(x_train)
check_null(x_test)
```

There are no missing values in the dataset  
 There are no missing values in the dataset

```
In [62]: # if for any columns , the variance is equal to zero, then you need to remove those variables.
# apply label encoder
for column in usable_columns:
    cardinality = len(np.unique(x_train[column]))
    if cardinality == 1:
        x_train.drop(column, axis= 1)
        x_test.drop(column,axis=1)
    if cardinality > 2:
        mapper = lambda x: sum ([ord(digit) for digit in x])
        x_train[column] = x_train[column].apply(mapper)
        x_test[column] = x_test[column].apply(mapper)
x_train.head()
```

/usr/local/lib/python3.7/site-packages/ipykernel\_launcher.py:8: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

/usr/local/lib/python3.7/site-packages/ipykernel\_launcher.py:9: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
if __name__ == '__main__':
```

Out[62]:

	X152	X196	X347	X70	X370	X8	X156	X263	X270	X248	...	X65	X68	X336	X204	X211
0	0	0	0	1	0	111	1	1	0	0	...	0	1	0	1	0
1	0	0	0	1	0	111	1	1	0	0	...	0	0	1	0	0
2	0	0	0	1	0	120	0	0	0	0	...	0	1	0	0	0
3	0	0	0	1	0	101	0	0	0	0	...	0	0	0	0	0
4	0	0	0	1	0	110	0	0	0	0	...	0	0	0	0	0

5 rows × 376 columns



```
In [63]: # check whether data is converted into numerical values or not.
print('Feature types:')
x_train[cols].dtypes.value_counts()
```

Feature types:

```
Out[63]: int64    376
dtype: int64
```

```
In [64]: # perform Dimensionality reduction
# linear dimensionality reduction using Singular Value Decomposition of Data to project it to Lower Dimensional Space.
n_comp = 12
pca = PCA(n_components = n_comp , random_state=420)
pca2_results_train = pca.fit_transform(x_train)
pca2_results_test = pca.fit_transform(x_test)
```

```
In [65]: # import the xgboost
# train the data using xgboost
import xgboost as xgb
from sklearn.metrics import r2_score
from sklearn.model_selection import train_test_split
```

```
In [66]: x_train, x_valid ,y_train, y_valid = train_test_split(
pca2_results_train,y_train ,test_size= 0.2,
random_state= 4242)
```

```
In [67]: d_train = xgb.DMatrix(x_train,label= y_train)
d_valid = xgb.DMatrix(x_valid,label= y_valid)

d_test = xgb.DMatrix(pca2_results_test)
```

```
In [68]: params = {}
params ['objective'] = 'reg:linear'
params ['eta'] = 0.02
params ['max_depth'] = 4
```

```
In [76]: def xgb_r2_score(preds , dtrain):
labels= dtrain.get_label()
return 'r2', r2_score(labels, preds)
```

```
In [77]: watchlist = [(d_train , 'train'), (d_valid, 'valid')]
```

```
In [78]: clf = xgb.train(params, d_train,  
                        1000, watchlist, early_stopping_rounds=50,  
                        feval=xgb_r2_score, maximize=True, verbose_eval=10)
```

[12:27:54] WARNING: /workspace/src/objective/regression\_obj.cu:167: reg:linear is now deprecated in favor of reg:squarederror.

[0] train-rmse:99.14835 valid-rmse:98.26297 train-r2:-58.35295 v  
alid-r2:-67.63754

Multiple eval metrics have been passed: 'valid-r2' will be used for early stopping.

Will train until valid-r2 hasn't improved in 50 rounds.

[10] train-rmse:81.27653 valid-rmse:80.36433 train-r2:-38.88428 v  
alid-r2:-44.91014

[20] train-rmse:66.71610 valid-rmse:65.77334 train-r2:-25.87403 v  
alid-r2:-29.75260

[30] train-rmse:54.86915 valid-rmse:53.89120 train-r2:-17.17724 v  
alid-r2:-19.64513

[40] train-rmse:45.24563 valid-rmse:44.22232 train-r2:-11.36018 v  
alid-r2:-12.90160

[50] train-rmse:37.44742 valid-rmse:36.37758 train-r2:-7.46672 v  
alid-r2:-8.40697

[60] train-rmse:31.15105 valid-rmse:30.01771 train-r2:-4.85891 v  
alid-r2:-5.40526

[70] train-rmse:26.08769 valid-rmse:24.90855 train-r2:-3.10906 v  
alid-r2:-3.41041

[80] train-rmse:22.04899 valid-rmse:20.82566 train-r2:-1.93528 v  
alid-r2:-2.08304

[90] train-rmse:18.84732 valid-rmse:17.59580 train-r2:-1.14472 v  
alid-r2:-1.20090

[100] train-rmse:16.33600 valid-rmse:15.07903 train-r2:-0.61125 v  
alid-r2:-0.61633

[110] train-rmse:14.40326 valid-rmse:13.14908 train-r2:-0.25254 v  
alid-r2:-0.22906

[120] train-rmse:12.93262 valid-rmse:11.69372 train-r2:-0.00982 v  
alid-r2:0.02795

[130] train-rmse:11.81574 valid-rmse:10.61241 train-r2:0.15707 v  
alid-r2:0.19941

[140] train-rmse:10.98584 valid-rmse:9.84577 train-r2:0.27132 v  
alid-r2:0.31090

[150] train-rmse:10.37818 valid-rmse:9.31608 train-r2:0.34970 v  
alid-r2:0.38305

[160] train-rmse:9.92761 valid-rmse:8.95044 train-r2:0.40494 v  
alid-r2:0.43053

[170] train-rmse:9.59297 valid-rmse:8.71236 train-r2:0.44438 v  
alid-r2:0.46042

[180] train-rmse:9.34889 valid-rmse:8.54634 train-r2:0.47229 v  
alid-r2:0.48079

[190] train-rmse:9.16216 valid-rmse:8.44332 train-r2:0.49316 v  
alid-r2:0.49323

[200] train-rmse:9.02020 valid-rmse:8.37881 train-r2:0.50875 v  
alid-r2:0.50095

[210] train-rmse:8.91339 valid-rmse:8.34352 train-r2:0.52031 v  
alid-r2:0.50514

[220] train-rmse:8.82193 valid-rmse:8.31729 train-r2:0.53011 v  
alid-r2:0.50825

[230] train-rmse:8.76377 valid-rmse:8.30382 train-r2:0.53628 v  
alid-r2:0.50984

[240] train-rmse:8.70541 valid-rmse:8.29665 train-r2:0.54244 v  
alid-r2:0.51069

[250] train-rmse:8.66769 valid-rmse:8.29059 train-r2:0.54639 v  
alid-r2:0.51140

[260] train-rmse:8.62392 valid-rmse:8.28817 train-r2:0.55097 v

alid-r2:0.51169			
[270] train-rmse:8.59049	valid-rmse:8.28617	train-r2:0.55444	v
alid-r2:0.51192			
[280] train-rmse:8.55208	valid-rmse:8.28102	train-r2:0.55841	v
alid-r2:0.51253			
[290] train-rmse:8.52721	valid-rmse:8.28241	train-r2:0.56098	v
alid-r2:0.51236			
[300] train-rmse:8.50348	valid-rmse:8.28322	train-r2:0.56342	v
alid-r2:0.51227			
[310] train-rmse:8.47840	valid-rmse:8.27920	train-r2:0.56599	v
alid-r2:0.51274			
[320] train-rmse:8.45433	valid-rmse:8.28000	train-r2:0.56845	v
alid-r2:0.51265			
[330] train-rmse:8.43036	valid-rmse:8.27771	train-r2:0.57090	v
alid-r2:0.51292			
[340] train-rmse:8.40901	valid-rmse:8.27708	train-r2:0.57307	v
alid-r2:0.51299			
[350] train-rmse:8.38952	valid-rmse:8.27791	train-r2:0.57504	v
alid-r2:0.51289			
[360] train-rmse:8.36299	valid-rmse:8.27466	train-r2:0.57773	v
alid-r2:0.51328			
[370] train-rmse:8.33332	valid-rmse:8.27140	train-r2:0.58072	v
alid-r2:0.51366			
[380] train-rmse:8.30903	valid-rmse:8.27299	train-r2:0.58316	v
alid-r2:0.51347			
[390] train-rmse:8.28632	valid-rmse:8.27065	train-r2:0.58543	v
alid-r2:0.51375			
[400] train-rmse:8.25956	valid-rmse:8.26821	train-r2:0.58811	v
alid-r2:0.51404			
[410] train-rmse:8.23725	valid-rmse:8.26749	train-r2:0.59033	v
alid-r2:0.51412			
[420] train-rmse:8.20771	valid-rmse:8.26607	train-r2:0.59326	v
alid-r2:0.51429			
[430] train-rmse:8.18372	valid-rmse:8.26409	train-r2:0.59564	v
alid-r2:0.51452			
[440] train-rmse:8.16154	valid-rmse:8.26225	train-r2:0.59783	v
alid-r2:0.51474			
[450] train-rmse:8.13499	valid-rmse:8.26102	train-r2:0.60044	v
alid-r2:0.51488			
[460] train-rmse:8.11711	valid-rmse:8.25870	train-r2:0.60219	v
alid-r2:0.51515			
[470] train-rmse:8.08796	valid-rmse:8.25732	train-r2:0.60504	v
alid-r2:0.51531			
[480] train-rmse:8.05856	valid-rmse:8.26027	train-r2:0.60791	v
alid-r2:0.51497			
[490] train-rmse:8.03767	valid-rmse:8.25967	train-r2:0.60994	v
alid-r2:0.51504			
[500] train-rmse:8.01613	valid-rmse:8.25948	train-r2:0.61203	v
alid-r2:0.51506			
[510] train-rmse:7.99323	valid-rmse:8.25895	train-r2:0.61424	v
alid-r2:0.51512			
[520] train-rmse:7.97607	valid-rmse:8.25734	train-r2:0.61589	v
alid-r2:0.51531			
Stopping. Best iteration:			
[472] train-rmse:8.08025	valid-rmse:8.25677	train-r2:0.60580	v
alid-r2:0.51538			



```
In [79]: # predict the test_df values using xgboost
p_test= clf.predict(d_test)

sub = pd.DataFrame()
sub['ID'] = id_test
sub['y'] = p_test
sub.to_csv('xgb.csv', index = False)

sub.head()
```

Out[79]:

	ID	y
0	1	83.886436
1	2	104.960732
2	3	83.411240
3	4	77.086838
4	5	97.411743